


Convolutional Neural Network

HPC-Gateway Fall School

Israt Jahan Tulin, Scientific Researcher, AI Consultant Team, HZDR, Germany
Date: November 04, 2025


CNN Architecture



https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html

Fully Connected Network vs. CNN Idea


$$\text{Total Parameters} = (16 \times 4 + 4) + (4 \times 1 + 1) = 73$$




3

Local Connectivity Concept

$$\text{Total Parameters} = (4 \times 4 + 4) + (4 \times 1 + 1) = 25$$




Local Connectivity Concept




NB: only “green” connections shown for simplicity.

Filter Representation – introducing weight filter




Example 1: Convolution Operation



Example 2: Convolution Operation


Input image




Input	Output						
3 ₀	3 ₁	2 ₂	1	0	12	12	17
0 ₂	0 ₂	1 ₀	3	1	10	17	19
3 ₀	1 ₁	2 ₂	2	3	9	6	14
2	0	0	2	2			
2	0	0	0	1			

<https://www.ibm.com/think/topics/convolutional-neural-networks>


Multiple Filters / Building CNNs



RGB image - 3 color channel



Convolution Operation for color image




Kernels Are Only Applied Where They Fully Fit on the Input

Input					Output		
3 ₀	3 ₁	2 ₂	1	0			
0 ₂	0 ₂	1 ₀	3	1			
3 ₀	1 ₁	2 ₂	2	3			
2	0	0	2	2			
2	0	0	0	1			

12	12	17
10	17	19
9	6	14

- We usually use odd numbers for filters so that they are applied symmetrically around our input data
- Notice in the GIF earlier that the output from applying our kernel was smaller than the input.

Padding: Ensures edge pixels are included in convolution




Padding: valid vs same

Padding: “valid”

3	5	2	7
4	1	3	8
6	3	8	2
9	6	1	5

1	2	1
2	1	2
1	1	2




55	52
57	50

Padding: “same”

output size to stay the same: use `padding = (kernel_size - 1) / 2`


0	0	0	0	0	0
0	3	5	2	7	0
0	4	1	3	8	0
0	6	3	8	2	0
0	9	6	1	5	0
0	0	0	0	0	0

1	2	1
2	1	2
1	1	2



19	26	46	22
29	55	52	40
42	57	50	43
36	46	44	19

Stride: Controls How Far the Kernel Moves Over the Input Image



Using of both strides and padding

- By default, our kernels are only applied where the filter fully fits on top of the input
- But we can control this behaviour and the size of our output with:
 - **padding**: “pads” the outside of the input with 0’s to allow the kernel to reach the boundary pixels
 - **strides**: controls how far the kernel “steps” over pixels.


Right side there's an example with:

- **padding=1**: we have 1 layer of 0’s around our border
- **strides=(2, 2)**: our kernel moves 2 data points to the right for each row, then moves 2 data points down to the next row

Input	Output
0 ₂ 0 ₀ 0 ₁ 0 0 0 0	1 6 5
0 ₁ 2 ₀ 2 ₀ 3 3 3 0	7 10 9
0 ₀ 0 ₁ 1 ₁ 3 0 3 0	7 10 8
0 2 3 0 1 3 0	
0 3 3 2 1 2 0	
0 3 3 0 2 3 0	
0 0 0 0 0 0 0	

Exercise

A two-dimensional, 3x3 convolutional filter is applied to a two-dimensional 4x4 input feature map (no padding added):



What is the shape of the output feature map?

- (A) 4 by 4
- (B) 3 by 3
- (C) 2 by 2

Calculation of output feature map


If the $m \times m$ image convolved with the $n \times n$ kernel with stride s and padding p ,

$$\text{Size of output image: } \left\lfloor \frac{m + (2 \times p) - n}{s} + 1 \right\rfloor \times \left\lfloor \frac{m + (2 \times p) - n}{s} + 1 \right\rfloor$$


Example: 4×4 image convolved with 3×3 kernel with stride and padding 1,

$$\text{Size of output image: } \left\lfloor \frac{4 + (2 \times 1) - 3}{1} + 1 \right\rfloor \times \left\lfloor \frac{4 + (2 \times 1) - 3}{1} + 1 \right\rfloor = 4 \times 4$$


Convolution for color image




Convolution for color image





Multiple filters for each color channel



Convolve with 3D filters



Effect of a Kernel on an Image



We can blur this image by applying a filter with the following weights:

$$\begin{bmatrix} 0.0625 & 0.125 & 0.0625 \\ 0.125 & 0.25 & 0.125 \\ 0.0625 & 0.125 & 0.0625 \end{bmatrix}$$

```
kernel = torch.tensor([[[[ 0.0625,  0.1250,  0.0625],  
[ 0.1250,  0.2500,  0.1250],  
[ 0.0625,  0.1250,  0.0625]]]])
```


```
plot_conv(image, kernel)
```

Effect of a Kernel on an Image

Original



Filtered



How about this one:

$$\begin{bmatrix} -2 & -1 & 0 \\ -1 & 1 & 1 \\ 0 & 1 & 2 \end{bmatrix}$$


```
kernel = torch.tensor([[[[-2, -1, 0],  
                         [-1, 1, 1],  
                         [0, 1, 2]]]])  
plot_conv(image, kernel)
```

Effect of a Kernel on an Image

Original



Filtered



[Here's a great website](#) where we can play around with other filters.

One more:

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

```
kernel = torch.tensor([[[[-1, -1, -1],  
                         [-1, 8, -1],  
                         [-1, -1, -1]]]])  
  
plot_conv(image, kernel)
```

Convolutional Layers

In PyTorch: convolutional layers are defined as `torch.nn.Conv2d`.

Key arguments to define a convolutional layer:


- a) `in_channels` → Number of input features (e.g.,
 - i) 1 for grayscale images
 - ii) 3 for RGB color images)
- b) `out_channels`: → Number of kernels (filters) to learn
 - i) Similar to the number of hidden neurons in a dense layer
- c) `kernel_size` → Size of the filter (e.g., 3×3, 5×5, 7×7)
- d) `Kernels (Filters)` → Kernels Are Weights That CNNs Learn during training to Recognize Patterns
- e) `stride`→ Step size of the filter as it moves across the image
- f) `padding`→ Extra pixels added around the border.
 - i) Ensures edge pixels are included in convolution

Example 1: Output with 1 Filter

https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html

```
conv_layer = torch.nn.Conv2d(1, 1, kernel_size=(3, 3))  
plot_convs(image, conv_layer)
```

Original



Filter 1



Example 2: Output with 2 Filters

https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html

```
conv_layer = torch.nn.Conv2d(1, 2, kernel_size=(3, 3))  
plot_convs(image, conv_layer)
```

Original



Filter 1



Filter 2



Example 3: Output with 3 Filters

https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html

```
conv_layer = torch.nn.Conv2d(1, 3, kernel_size=(5, 5))  
plot_convs(image, conv_layer)
```

Original



Filter 1



Filter 2




Filter 3



Example 4: Effect of No Padding



```
conv_layer = torch.nn.Conv2d(1, 1, kernel_size=(51, 51))  
plot_convs(image, conv_layer, axis=True)
```



- Output image becomes smaller
- Edges lost since kernel doesn't cover full image
- Using a larger kernel shows this effect clearly

Example 5: With Padding



```
conv_layer = torch.nn.Conv2d(1, 1, kernel_size=(51, 51), padding=25)  
plot_convs(image, conv_layer, axis=True)
```



- **Padding** is added to the outside of the image to avoid this
- Setting **padding = kernel_size // 2** will always generate an output of the same shape as the input.

Example 6: With Strides

```
conv_layer = torch.nn.Conv2d(1, 1, kernel_size=(25, 25), stride=3)  
plot_convs(image, conv_layer, axis=True)
```





strides influence the size of the output

Features → Output Channels


Original Image

1 Input Channels

Convolutional Layer
5 Output Channels / Kernels



Features → Output Channels



34

https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html

Let's make the simple CNN in PyTorch

20,000 parameters in that last layer, that's a lot of parameters

Is there a way we can reduce this?

please see appendix for more better understanding

35

```
class CNN(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.main = torch.nn.Sequential(
            torch.nn.Conv2d(in_channels=1, out_channels=3, kernel_size=(3, 3), padding=1),
            torch.nn.ReLU(),
            torch.nn.Conv2d(in_channels=3, out_channels=2, kernel_size=(3, 3), padding=1),
            torch.nn.ReLU(),
            torch.nn.Flatten(),
            torch.nn.Linear(20000, 1)
        )

    def forward(self, x):
        out = self.main(x)
        return out
```

```
model = CNN()
summary(model, (1, 100, 100));
```

Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 1]	--
└─Conv2d: 2-1	[-1, 3, 100, 100]	30
└─ReLU: 2-2	[-1, 3, 100, 100]	--
└─Conv2d: 2-3	[-1, 2, 100, 100]	56
└─ReLU: 2-4	[-1, 2, 100, 100]	--
└─Flatten: 2-5	[-1, 20000]	--
└─Linear: 2-6	[-1, 1]	20,001


```
Total params: 20,087
Trainable params: 20,087
Non-trainable params: 0
Total mult-adds (M): 0.85
```

```
Input size (MB): 0.04
Forward/backward pass size (MB): 0.38
Params size (MB): 0.08
Estimated Total Size (MB): 0.50
```

Pooling Layers in CNN

- ❑ Used to **reduce dimensionality** of feature maps
→ fewer parameters after flattening (via `torch.nn.Flatten()`)
- ❑ Common pooling methods:
 - **Max Pooling**
 - **Average Pooling**
- ❑ **Max Pooling** often performs better
→ captures the **sharpest and most important features** of the image

Example: applying **Max Pooling** to highlight key patterns in a transformed image



https://www.researchgate.net/publication/333593451_Application_of_Transfer_Learning_Using_Convolutional_Neural_Network_Method_for_Early_Detection_of_Terry's_Nail

Implementing pooling with `torch.nn.MaxPool2d()`

*reduce the number
of parameters*

37

We reduced that last layer to 1,251 parameters
(~16 times smaller than the original number)

© Israt Ja

```
class CNN(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.main = torch.nn.Sequential(
            torch.nn.Conv2d(in_channels=1, out_channels=3, kernel_size=(3, 3), padding=1),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d((2, 2)),
            torch.nn.Conv2d(in_channels=3, out_channels=2, kernel_size=(3, 3), padding=1),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d((2, 2)),
            torch.nn.Flatten(),
            torch.nn.Linear(1250, 1)
        )

    def forward(self, x):
        out = self.main(x)
        return out
```


```
model = CNN()
summary(model, (1, 100, 100));
```

Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 1]	--
└ Conv2d: 2-1	[-1, 3, 100, 100]	30
└ ReLU: 2-2	[-1, 3, 100, 100]	--
└ MaxPool2d: 2-3	[-1, 3, 50, 50]	--
└ Conv2d: 2-4	[-1, 2, 50, 50]	56
└ ReLU: 2-5	[-1, 2, 50, 50]	--
└ MaxPool2d: 2-6	[-1, 2, 25, 25]	--
└ Flatten: 2-7	[-1, 1250]	--
└ Linear: 2-8	[-1, 1]	1,251


```
Total params: 1,337
Trainable params: 1,337
Non-trainable params: 0
Total mult-adds (M): 0.41
```

```
Input size (MB): 0.04
Forward/backward pass size (MB): 0.27
Params size (MB): 0.01
Estimated Total Size (MB): 0.31
```

Example: CNN Architecture




AlexNet Architecture Overview



First CNN to use GPU for faster training performance

- **5 Convolutional layers** → extract image features
- **3 Max-Pooling layers** → reduce spatial dimensions
- **2 Fully Connected layers** → combine learned features
- **1 Softmax layer** → output final class probabilities
- Each Conv layer uses **ReLU activation** (introduces non-linearity)
- **Input size:** $\sim 227 \times 227 \times 3$ (often referred to as $224 \times 224 \times 3$)

VGG-Net Architecture



40

<https://medium.com/latinixinai/convolutional-neural-network-from-scratch-6b1c856e1c07>



Thank You Any Questions?

Time for exercise



Appendix: Convolution using pytorch for gray scale image

https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html


```
nn.Conv2d(  
    in_channels=1,  
    out_channels=3,  
    kernel_size=(3,3)  
)
```

Total number
of parameters:

- 1 input channel
- 3 output channels
- (3×3) filters

$$1 \times 3 \times (3 \times 3) = 27 \text{ parameters}$$

$$+ 3 \text{ biases} = 30 \text{ parameters}$$



1 Input
Channel

3 Output
Channels

Appendix: Convolution using pytorch for gray scale image

https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html


```
nn.Conv2d(  
    in_channels=3,  
    out_channels=3,  
    kernel_size=(3,3)  
)
```

Total number
of parameters:

- 3 input channels
- 3 output channels
- (3 x 3) filters

$$3 \times 3 \times (3 \times 3)
= 81 \text{ parameters}$$

$$+ 3 \text{ biases}
= 84 \text{ parameters}$$



Appendix: Understanding Tensor Dimensions in PyTorch

PyTorch needs data in a **very specific shape** for convolution operations.

Image Tensor Shape

All images (inputs or outputs of convolution) are stored as **4D tensors** in PyTorch, with this shape:

(batch_size,n_channels,height,width)

So for one grayscale image of 100×100 , shape should be:

[1, 1, 100, 100]

Dimension	Meaning	Example
batch_size	Number of images processed together	1 (if just one image)
n_channels	Number of color channels	1 (grayscale) or 3 (RGB)
height	Image height (pixels)	100
width	Image width (pixels)	100

Appendix: Kernel (Filter) Tensor Shape

Each convolutional layer also has a **4D tensor** for its filters:

(n_kernels, n_channels, kernel_height, kernel_width)

Example: `conv_layer = torch.nn.Conv2d(1, 1, kernel_size=(5, 5))`

means:

- 1 input channel (grayscale)
- 1 output channel (so 1 kernel)
- Each kernel is 5×5 in size

Dimension	Meaning
n_kernels	Number of filters (output feature maps)
n_channels	Must match number of input channels
kernel_height, kernel_width	Filter size (e.g., 3×3, 5×5)

Appendix: Why the Shape Error Happened

First loaded the image:

```
image.shape  
torch.Size([100, 100, 4])
```

That's the **NumPy-style format** — (height, width, channels) (note the channel last).

But PyTorch expects **channels first**, like (channels, height, width) — and with an extra dimension for batch.

So we have to reshape it like this:

```
image = image[None, None, :]  
image.shape  
torch.Size([1, 1, 100, 100])
```

This means:

- 1 batch
- 1 channel
- 100×100 pixels

Now PyTorch knows exactly how to apply the convolution.

Appendix: Why Output Is Smaller

- After convolution:

```
conv_out.shape  
torch.Size([1, 1, 96, 96])
```

- It shrinks because the kernel (5×5) can't fit around edge pixels — this is “**no padding**”.
- Formula for output size (*with stride=1, no padding*):

$$\text{Output size} = \text{Input size} - \text{Kernel size} + 1$$

So:

$$100 - 5 + 1 = 96$$

Appendix: Visualizing the Result

To plot it with Matplotlib, a few clean-up steps are needed:

1. **Detach:**

`.detach()` removes the tensor from PyTorch's *computational graph* (so no gradients are tracked).

2. **Squeeze:**

`.squeeze()` removes any dimensions of size 1 (batch and channel), leaving only the 2D image.

3. **Plot:**


Then `plt.imshow()` can be used to show the image.

Example:

```
plt.imshow(conv_out.detach().squeeze(), cmap='gray')
```

Now the shape is [96, 96] — which Matplotlib understands.

48



https://ubc-mds.github.io/DSCI_572_sup-learn-2/lectures/06_cnns-pt1.html

Appendix: Flattening

- With our convolutional layers, we're basically just passing images through the network
- But we're going to eventually want to do some regression or classification
- That means that by the end of our network, we are going to need to `torch.nn.Flatten()` our images:

